

Introduction

This report presents a comprehensive analysis of a dataset containing details about used cars, including attributes such as manufacturing year, selling price, kilometers driven, fuel type, seller type, transmission type, ownership history, mileage, engine capacity, maximum power, and seating capacity. The objective is to explore and understand the factors influencing the selling price of used cars. Through data cleaning, transformation, and exploratory data analysis, key trends and correlations are identified, providing insights into the used car market. This analysis serves as a foundation for predictive modeling and further statistical analysis to better understand price determinants in the used car market.

Importing Libraries

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from datetime import datetime
import numpy as np
from sklearn.feature_selection import SelectKBest, f_regression
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
```

Dataset Used

df

| | name | year | selling_price | km_driven | fuel | seller_type | transmission | owner | mileage | engine | max_power | torque | seats |
|---|------------------------------|------|---------------|-----------|--------|-------------|--------------|--------------|------------|---------|------------|--------------------------|-------|
| 0 | Maruti Swift Dzire VDI | 2014 | 450000 | 145500 | Diesel | Individual | Manual | First Owner | 23.4 kmpl | 1248 CC | 74 bhp | 190Nm@ 2000rpm | 5.0 |
| 1 | Skoda Rapid 1.5 TDI Ambition | 2014 | 370000 | 120000 | Diesel | Individual | Manual | Second Owner | 21.14 kmpl | 1498 CC | 103.52 bhp | 250Nm@ 1500-2500rpm | 5.0 |
| 2 | Honda City 2017-2020 EXi | 2006 | 158000 | 140000 | Petrol | Individual | Manual | Third Owner | 17.7 kmpl | 1497 CC | 78 bhp | 12.7@ 2,700(kgm@ rpm) | 5.0 |
| 3 | Hyundai i20 Sportz Diesel | 2010 | 225000 | 127000 | Diesel | Individual | Manual | First Owner | 23.0 kmpl | 1396 CC | 90 bhp | 22.4 kgm at 1750-2750rpm | 5.0 |
| 4 | Maruti Swift VXi BSIII | 2007 | 130000 | 120000 | Petrol | Individual | Manual | First Owner | 16.1 kmpl | 1298 CC | 88.2 bhp | 11.5@ 4,500(kgm@ rpm) | 5.0 |

Data Overview

The dataset contains information about used cars, with attributes such as the car's name, manufacturing year, selling price, kilometers driven, fuel type, seller type, transmission type, ownership history, mileage, engine capacity, maximum power, and number of seats.

Data Cleaning and Preparation

- **Removing Null Values:** We identified and removed null values present in the mileage, engine, max power, and seats columns.
- **Handling Duplicates:** We detected and eliminated duplicate records to ensure the dataset's integrity.

```
df.drop(columns=['torque'],inplace = True)
#checking null values
df.isnull().sum()
```

```
name          0
year          0
selling_price 0
km_driven     0
fuel          0
seller_type   0
transmission  0
owner         0
mileage       221
engine        221
max_power     215
seats         221
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8128 entries, 0 to 8127
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   name                8128 non-null   object
1   year                8128 non-null   int64
2   selling_price       8128 non-null   int64
3   km_driven           8128 non-null   int64
4   fuel                8128 non-null   object
5   seller_type         8128 non-null   object
6   transmission        8128 non-null   object
7   owner               8128 non-null   object
8   mileage             7907 non-null   object
9   engine              7907 non-null   object
10  max_power           7913 non-null   object
11  seats               7907 non-null   float64
dtypes: float64(1), int64(3), object(8)
memory usage: 762.1+ KB
```

```
#removing null values
df.dropna(inplace=True)
```

- **Extracting Useful Features:**

- **No. of Years Used:** A new column, “no_year”, was created to represent the number of years a car has been used, calculated by subtracting the manufacturing year from the current year (2024).

- **Encoding Categorical Variables:**

- Fuel type, seller type, transmission type, owner type, and car brand names were encoded as numerical values for better analysis and model training.

```
#creating a new column No_year(Number of years car used)
```

```
current_datetime=datetime.now()
current_year=current_datetime.year
print(current_year)
df["no_year"]=current_year-df["year"]
2024
```

```
# Drop column Year
```

```
df.drop(['year'],axis=1,inplace=True)
```

```
# Fuel_Type column : 'Diesel':1, 'Petrol':2, 'LPG':3, 'CNG':4
```

```
print(df['fuel'].unique())
df['fuel'].replace({'Diesel':1, 'Petrol':2, 'LPG':3, 'CNG':4},inplace=True)
['Diesel' 'Petrol' 'LPG' 'CNG']
```

```
# Selling_type column : 'Individual':1, 'Dealer':2, 'Trustmark Dealer':3
```

```
print(df['seller_type'].unique())
df['seller_type'].replace({'Individual':1, 'Dealer':2, 'Trustmark Dealer':3},inplace=True)
['Individual' 'Dealer' 'Trustmark Dealer']
```

```
# Transmission column :replacing manual with 1 and automatic with 2
```

```
print(df['transmission'].unique())
df['transmission'].replace({'Manual':1, 'Automatic':2},inplace=True)
['Manual' 'Automatic']
```

```
#OWNER column: 'First Owner':1, 'Second Owner':2, 'Third Owner':3, 'Fourth & Above Owner':4, 'Test Drive Car':5
```

```
print(df['owner'].unique())
df['owner'].replace({'First Owner':1, 'Second Owner':2, 'Third Owner':3, 'Fourth & Above Owner':4, 'Test Drive Car':5},inplace=True)
['First Owner' 'Second Owner' 'Third Owner' 'Fourth & Above Owner'
 'Test Drive Car']
```

Data Transformation

- **Modifying Mileage, Engine, and Max Power:** These columns originally contained text values with units. We extracted the numerical part of these values for analysis.
- **Brand Name Extraction:** The car brand was extracted from the name column and encoded into numerical values.

function to extract value from mileage, engine and max_power

```
def clean_data(value):
    value=str(value)
    value = value.split(' ')[0]
    value = value.strip()
    if value == '':
        return np.nan
    else:
        return float(value)
```

```
df['mileage'] = df['mileage'].apply(clean_data)
```

```
df['engine'] = df['engine'].apply(clean_data)
```

```
df['max_power'] = df['max_power'].apply(clean_data)
```

```
df.head(3)
```

| | name | selling_price | km_driven | fuel | seller_type | transmission | owner | mileage | engine | max_power | seats | no_year |
|---|------|---------------|-----------|------|-------------|--------------|-------|---------|--------|-----------|-------|---------|
| 0 | 1 | 450000 | 145500 | 1 | 1 | 1 | 1 | 23.40 | 1248.0 | 74.00 | 5.0 | 10 |
| 1 | 2 | 370000 | 120000 | 1 | 1 | 1 | 2 | 21.14 | 1498.0 | 103.52 | 5.0 | 10 |
| 2 | 3 | 158000 | 140000 | 2 | 1 | 1 | 3 | 17.70 | 1497.0 | 78.00 | 5.0 | 18 |

Exploratory Data Analysis

- **Distributions:** We analyzed the distribution of key numerical features (selling price, kilometers driven, mileage, engine capacity, max power, and number of years used) using distribution plots.
 - **Selling Price:** The distribution is right-skewed, indicating that most cars are sold at a lower price.
 - **Kilometers Driven:** This distribution is also right-skewed, with most cars having fewer kilometers driven.
 - **Mileage:** There is a wide range in mileage, but a significant number of cars offer moderate mileage.
 - **Engine Capacity:** The distribution shows a concentration of cars with smaller engine sizes.
 - **Max Power:** Most cars have lower maximum power values.
 - **Number of Years Used:** The distribution shows a mix of newer and older cars.
 - **Box plots:** Box plots were used to identify outliers in the dataset. Significant outliers were present in the selling price, kilometers driven, mileage, engine capacity, and max power columns.

Correlation Analysis - A correlation matrix was generated to understand the relationships between numerical variables.

- **Positive Correlations:**
 - Selling price is positively correlated with engine capacity and max power.
 - Engine capacity and max power also show a strong positive correlation.
- **Negative Correlations:**
 - Selling price is negatively correlated with the number of years used, indicating that older cars tend to sell for less. Mileage has a negative correlation with engine capacity and max power.

BEST FEATURES

```
bestfeatures = SelectKBest(score_func=f_regression, k=10)

# Fit SelectKBest to your data
fit = bestfeatures.fit(x, y)

# Create DataFrame to store feature scores
dfscores = pd.DataFrame(fit.scores_)

# Create DataFrame to store feature names
dfcolumns = pd.DataFrame(x.columns)

# Concatenate feature names and scores into a single DataFrame for better visualization
featureScores = pd.concat([dfcolumns, dfscores], axis=1)

# Rename the columns
featureScores.columns = ['Specs', 'Score']

# Print the top 10 features with the highest scores
print(featureScores.nlargest(10, 'Score'))
```

| | Specs | Score |
|----|--------------|-------------|
| 8 | max_power | 6181.382158 |
| 4 | transmission | 1857.986653 |
| 7 | engine | 1637.476893 |
| 10 | no_year | 1500.224991 |
| 0 | name | 453.852706 |
| 3 | seller_type | 449.284164 |
| 2 | fuel | 430.886346 |
| 5 | owner | 307.528682 |
| 1 | km_driven | 179.295418 |
| 9 | seats | 173.112006 |

Correlation Matrix

```
correlation_matrix = df.corr()
correlation_matrix
```

| | name | selling_price | km_driven | fuel | seller_type | transmission | owner | mileage | engine | max_power | seats | no_year |
|---------------|-----------|---------------|-----------|-----------|-------------|--------------|-----------|-----------|-----------|-----------|-----------|-----------|
| name | 1.000000 | 0.251613 | 0.049784 | -0.277496 | 0.047916 | 0.174458 | -0.013379 | -0.273961 | 0.369460 | 0.348463 | 0.155834 | -0.033978 |
| selling_price | 0.251613 | 1.000000 | -0.161265 | -0.245558 | 0.250423 | 0.465538 | -0.209265 | -0.108655 | 0.442772 | 0.692323 | 0.158531 | -0.427335 |
| km_driven | 0.049784 | -0.161265 | 1.000000 | -0.252491 | -0.126336 | -0.118965 | 0.252205 | -0.196419 | 0.253460 | 0.041770 | 0.207890 | 0.387918 |
| fuel | -0.277496 | -0.245558 | -0.252491 | 1.000000 | -0.019724 | 0.005210 | -0.012138 | -0.035961 | -0.510633 | -0.328039 | -0.343668 | 0.043564 |
| seller_type | 0.047916 | 0.250423 | -0.126336 | -0.019724 | 1.000000 | 0.213725 | -0.151667 | 0.001552 | 0.065629 | 0.187339 | -0.040726 | -0.148137 |
| transmission | 0.174458 | 0.465538 | -0.118965 | 0.005210 | 0.213725 | 1.000000 | -0.076854 | -0.173667 | 0.219526 | 0.441681 | -0.019314 | -0.143997 |
| owner | -0.013379 | -0.209265 | 0.252205 | -0.012138 | -0.151667 | -0.076854 | 1.000000 | -0.188624 | 0.033741 | -0.052018 | 0.007649 | 0.480096 |
| mileage | -0.273961 | -0.108655 | -0.196419 | -0.035961 | 0.001552 | -0.173667 | -0.188624 | 1.000000 | -0.579153 | -0.378609 | -0.459188 | -0.366048 |
| engine | 0.369460 | 0.442772 | 0.253460 | -0.510633 | 0.065629 | 0.219526 | 0.033741 | -0.579153 | 1.000000 | 0.683506 | 0.658711 | 0.019763 |
| max_power | 0.348463 | 0.692323 | 0.041770 | -0.328039 | 0.187339 | 0.441681 | -0.052018 | -0.378609 | 0.683506 | 1.000000 | 0.259028 | -0.159889 |
| seats | 0.155834 | 0.158531 | 0.207890 | -0.343668 | -0.040726 | -0.019314 | 0.007649 | -0.459188 | 0.658711 | 0.259028 | 1.000000 | -0.025021 |
| no_year | -0.033978 | -0.427335 | 0.387918 | 0.043564 | -0.148137 | -0.143997 | 0.480096 | -0.366048 | 0.019763 | -0.159889 | -0.025021 | 1.000000 |

Insights

- **Factors Influencing Car Prices:** Engine capacity, max power, and the age of the car are significant factors affecting the selling price. Cars with higher engine capacity and max power tend to sell for more, while older cars generally fetch lower prices.
- **Model Training:** Implemented Random Forest regression model to predict the selling price and evaluate their performance using metrics like Mean Squared Error (MSE), Mean Absolute Error(MAE) and R-squared (R^2) values.

Random Forest Regressor Model

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

# Split the data into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y_scaled, test_size=0.2, random_state=42)

# Create and train the Random Forest Regression model
rf_model = RandomForestRegressor(n_estimators=10, random_state=45)
rf_model.fit(x_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(x_test)

# Evaluate the model
mae=metrics.mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"MAE score:{mae}")
print(f"Mean Squared Error: {mse}")
print(f"R^2 Score: {r2}")
```

```
MAE score:0.1530463136374124
Mean Squared Error: 0.08372642537719306
R^2 Score: 0.8954275160815178
```

MODEL INPUT

```
: columns = ['name', 'km_driven', 'fuel', 'seller_type', 'transmission', 'owner', 'mileage', 'engine', 'max_power', 'seats', 'no_year']
dat = pd.DataFrame([[4, 90000, 1, 1, 1, 19.67, 1582.0, 126.20, 5.0, 7]], columns=columns)
```

SCALE THE INPUT DATA AND INVERSEING THE OUTPUT BACK TO ORIGINAL SCALE

```
: # Ensure the new data for prediction has the same columns as x
dat = dat[x.columns]

# Scale the new data
dat_scaled = scaler_x.transform(dat)

# Make predictions using the model
predictions_scaled = rf_model.predict(dat_scaled)

# Inverse transform the predictions to get them back to the original scale
predictions = scaler_y.inverse_transform(predictions_scaled.reshape(-1, 1))

# Print the predictions in the original scale
print("Predictions (original scale):", predictions)

Predictions (original scale): [[1155399.9]]
```